

Learning to Act on Screens: VLM-RL for Real-World GUI Automation

Fang Sun

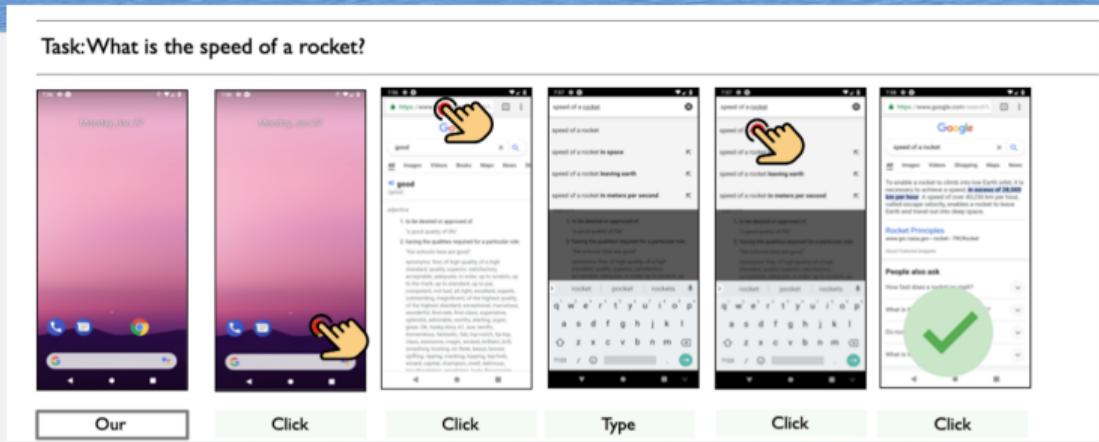
Nov 7, 2025



Project Summary

- ▶ Led an end-to-end effort to train a **1.3B VLM** as a **general-purpose GUI agent**
- ▶ Built from **700,000 Android GUI interaction trajectories** (offline dataset)
- ▶ Used [vLLM](#) for high-throughput inference and [DeepSpeed ZeRO-3](#) for memory-efficient fine-tuning
- ▶ Designed an **Offline-PPO** algorithm with [stochastic-aware advantage estimation](#) to handle GUI non-stationarity and interface variability
- ▶ Achieved **53.8% task success**, outperforming:
 - ▶ Supervised baseline: **+29.5% absolute**
 - ▶ GPT-4V: **8.3%**
 - ▶ CogAgent: **38.5%**

How the Agent Works: From Perception to Action



- ▶ Task: "What is the speed of a rocket?"
- ▶ The agent observes the **entire screen** at each step.
- ▶ Selects actions sequentially: **Click** → **Click** → **Type** → **Click** → **Click**.
- ▶ ~~Learns why actions move the task forward, not where buttons are.~~
- ▶ Enables **generalization** across UI layouts and devices.

Initiative: What Sparked This Project

- ▶ Many internal workflows (reporting, form submission, testing) relied on **rule-based GUI scripts**.
- ▶ These scripts were **fragile**: small UI layout changes caused repeated failures.
- ▶ Even advanced **VLMs (e.g., GPT-4V)** could **see** the interface but **could not decide the next action**.
- ▶ Early supervised fine-tuning attempts **overfit to pixel positions** instead of learning task strategies.
- ▶ I recognized the core challenge was not perception, but **decision-making under UI variability**.

Initiative: What I Did and The Immediate Impact

- ▶ I reframed GUI automation as a **sequential decision problem** and proposed using **reinforcement learning**.
- ▶ Designed a **safe offline RL pipeline** to avoid real-time interaction risks.
- ▶ Initiated and coordinated collection of **700k Android GUI interaction trajectories**.
- ▶ Led a staged training plan: **imitation learning** → **advantage-weighted offline RL**.
- ▶ **Direct Result:** The team shifted from **script-based automation** to a **learning-based agentic framework**, enabling models that **act**, not just **see**.

Innovation: Offline-PPO for GUI Action Learning

- ▶ We cannot safely or efficiently **explore online** in GUI environments:
 - ▶ Wrong clicks can lose progress, break workflows, or corrupt state
 - ▶ Real systems do not allow trial-and-error at high frequency
- ▶ We use an **advantage-weighted behavioral cloning** objective to learn action preferences:

$$\max_{\pi} \mathbb{E}_{(s,a) \sim D} \left[\exp\left(\frac{A(s,a)}{\beta}\right) \log \pi(a|s) \right] - \lambda \cdot \text{KL}(\pi \parallel \pi_{\text{behavior}})$$

- ▶ Interpretation:
 - ▶ Good actions (positive advantage) get **higher probability**
 - ▶ Bad / irrelevant actions are **down-weighted**
 - ▶ KL regularization prevents the policy from drifting away from **human-like behavior**
- ▶ Result: We transform a passive VLM into a policy that chooses actions, without performing any online exploration.

Innovation: Stochastic-Aware Advantage Estimation

- ▶ GUI environments are **visually and functionally non-stationary**:
 - ▶ UI layout shifts with resolution or window scaling
 - ▶ Theme packs, OS versions, or app updates modify colors or widget shapes
 - ▶ Dynamic pop-ups introduce unexpected states
- ▶ Naive advantage estimation would **overfit** to specific screen appearances or trajectories.
- ▶ Our solution adds **stochastic smoothing and clipping**:
 - ▶ **Temporal smoothing** stabilizes advantages across similar states
 - ▶ **Clipping** prevents extremely large advantages from dominating learning
 - ▶ **State augmentation** increases tolerance to layout perturbation
- ▶ Effect:
 - ▶ The agent **generalizes across UI variations** rather than memorizing pixel layouts
 - ▶ The model becomes **robust to unseen screen states** and dynamic interface changes

Implementation Process & Collaboration

- ▶ I initiated and led the project end-to-end, from problem framing to prototypes to full training pipeline.
- ▶ **Solo-led core development:**
 - ▶ Built dataset pipeline (collection, filtering, trajectory formatting)
 - ▶ Implemented imitation learning and Offline-PPO training loop
 - ▶ Designed action-token interface and runtime execution layer
- ▶ **Collaborative inputs where needed:**
 - ▶ Advisor feedback on algorithmic strategy
 - ▶ Occasional engineering support for environment instrumentation
- ▶ Result: A **fully functional, reproducible pipeline** for training agentic GUI VLMs.

Implementation: Dataset Construction

- ▶ Collected **human demonstration trajectories** from real GUI workflows
- ▶ Each step contains:
 - ▶ Screenshot state s_t
 - ▶ Task instruction (goal)
 - ▶ Action token a_t (click / drag / type / confirm)
 - ▶ Reward signal r_t representing task progress
- ▶ Stored as transitions:
$$(s_t, a_t, r_t, s_{t+1})$$
- ▶ Filtered out:
 - ▶ Idle cursor motion
 - ▶ Accidental clicks
 - ▶ Dead-end trajectories
- ▶ Result: Clean dataset reflecting intentional expert behavior.

Implementation: Imitation Warm-Start & Offline RL Refinement

Step 1: Supervised Imitation (Behavior Cloning)

- ▶ Model learns to **imitate expert actions** directly from demonstrations
- ▶ Prevents unstable “random exploration” behavior

Step 2: Offline-PPO Policy Refinement

$$\max_{\pi} \mathbb{E}_{(s,a) \sim D} \left[\exp\left(\frac{A(s,a)}{\beta}\right) \log \pi(a|s) \right] - \lambda \text{KL}(\pi \parallel \pi_{\text{behavior}})$$

- ▶ **Increase** probability of actions that advance task progress
- ▶ **Reduce** probability of suboptimal actions
- ▶ **Constrain policy** near demonstrated behavior for stability & safety

Implementation: Action Execution Runtime & Training Infrastructure

Action Execution Layer

- ▶ Converts predicted tokens to executable actions:
- ▶ CLICK(x, y), SELECT, TYPE("text"), CONFIRM
- ▶ Decouples policy learning from UI rendering to **portable across GUIs**

Training Infrastructure

- ▶ [vLLM](#): High-throughput inference during evaluation loops
- ▶ [DeepSpeed ZeRO-3](#): Sharded training for the 1.3B parameter policy model

Outcome

The system trains efficiently on **commodity multi-GPU hardware** and outputs a **robust, multi-step GUI automation agent**.

Insights: Why Previous Approaches Struggled

- ▶ **Visual Non-Stationarity is the Core Challenge**
 - ▶ UI layouts change over time (position, size, color, pop-ups)
 - ▶ Supervised VLMs **memorize** training layouts → fail when screens shift
- ▶ **The Problem Was Not Recognition — It Was Decision Adaptation**
 - ▶ Models could see elements correctly
 - ▶ They did not know what to do next when context changed
- ▶ **Key Insight:** GUI automation must be framed as **policy decision-making**, not screen captioning.

Insights: Limitations of Imitation Learning

- ▶ **Imitation Learning Reproduces Patterns, Not Strategy**
 - ▶ Behavior cloning copies human actions seen during training
 - ▶ Breaks down in new or unseen states
- ▶ GUI tasks require **credit assignment across multiple steps**
 - ▶ E.g., clicking a menu now enables success several steps later
- ▶ **Offline Reinforcement Learning was required** to learn:
 - ▶ **Why** actions matter (not just what they look like)
 - ▶ How to recover from unfamiliar or shifted UI states
- ▶ **Key Insight:** GUI automation requires **policy reasoning**, not action mimicry.

Insights: What Enabled Generalization

- ▶ **Advantage-Weighted Offline PPO Improved Policy Quality**
 - ▶ Increased probability of actions that reliably advanced task progress
 - ▶ Worked entirely offline — no risky or costly real-time exploration
- ▶ **Decoupling Actions from Pixel Coordinates Was Critical**
 - ▶ Represented actions as semantic UI intents (e.g., “confirm”, “open menu”)
 - ▶ Dramatically reduced brittleness across applications and screen layouts
- ▶ **Key Insight:** Stable generalization comes from

(offline credit assignment) + (semantic action abstraction).

Iteration: Refining the Approach

- ▶ **Policy Drift During Offline RL**
 - ▶ Early policies deviated from human behavior to unstable “random clicking”
 - ▶ Solution: **KL regularization + advantage normalization**
 - ▶ Result: Stable updates that stay close to expert intent
- ▶ **Ambiguity in Multi-Element Screens**
 - ▶ Model misclicked when multiple similar buttons appeared (e.g., multiple “Next” buttons)
 - ▶ Solution: **Semantic action abstraction + goal-grounded UI prompts**
 - ▶ Result: Better element selection precision
- ▶ **Overfitting to Frequent UI Layouts**
 - ▶ Model performed well on common states, poorly on rare or unseen screens
 - ▶ Solution: **UI theme / resolution randomization + synthetic pop-ups**
 - ▶ Result: Improved generalization across interface variation

Future Improvements: Increasing Velocity & Efficiency

- ▶ **More Systematic Evaluation Signals**
 - ▶ Added mid-trajectory checkpoints to diagnose where reasoning failed
 - ▶ Going forward: Integrate **step-level success scores** automatically during data logging
- ▶ **Smarter Dataset Curation**
 - ▶ Next iteration would **actively mine failure states** for targeted offline fine-tuning
 - ▶ Reduces training time by focusing on high-impact cases
- ▶ **Faster Iteration Loops**
 - ▶ Early loops required full training runs to observe behavior changes
 - ▶ Would adopt **agent replay + on-device lightweight rollouts** for faster debugging
- ▶ **Scalable Action Abstraction Library**
 - ▶ Future improvement: A reusable **semantic UI action vocabulary** across apps
 - ▶ ~~Reduces retraining effort when transferring to new GUI environments~~

Impact

Quantitative Results

Model	Success
GPT-4V	8.3%
BC (Supervised VLM)	24.3%
CogAgent	38.5%
Ours (Offline-RL VLM)	53.8%

- ▶ +29.5% absolute over supervised
- ▶ 5–6x higher than GPT-4V prompting
- ▶ 1.4x over state-of-the-art agent

Qualitative Outcomes

- ▶ Generalizes to **unseen UI layouts**
- ▶ Robust to **pop-ups, themes, resizing**
- ▶ Previously manual workflows ■ **fully automated**

Key Insight

- ▶ GUI automation is a **policy decision problem**, not just visual recognition
- ▶ Offline RL enables safe and scalable improvement **without online trials**

Thank You for the opportunity!